



CART-RF-ANN



PREPARED BY
MURALIDHARAN N

CART-RF-ANN

An Insurance firm providing tour insurance is facing higher claim frequency. The management decides to collect data from the past few years. You are assigned the task to make a model which predicts the claim status and provide recommendations to management. Use CART, RF & ANN and compare the models' performances in train and test sets.

Data Dictionary

1. Target: Claim Status (Claimed)
2. Code of tour firm (Agency Code)
3. Type of tour insurance firms (Type)
4. Distribution channel of tour insurance agencies (Channel)
5. Name of the tour insurance products (Product)
6. Duration of the tour (Duration)
7. Destination of the tour (Destination)
8. Amount of sales of tour insurance policies (Sales)
9. The commission received for tour insurance firm (Commission)
10. Age of insured (Age)

2.1 Data Ingestion: Read the dataset. Do the descriptive statistics and do null value condition check, write an inference on it.

Importing all required Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, roc_curve, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

Reading the dataset,

Checking the data

```
df_insured.head()
```

[5]:

	Age	Agency_Code	Type	Claimed	Commision	Channel	Duration	Sales	Product Name	Destination
0	48	C2B	Airlines	No	0.70	Online	7	2.51	Customised Plan	ASIA
1	36	EPX	Travel Agency	No	0.00	Online	34	20.00	Customised Plan	ASIA
2	39	CWT	Travel Agency	No	5.94	Online	3	9.90	Customised Plan	Americas
3	36	EPX	Travel Agency	No	0.00	Online	4	26.00	Cancellation Plan	ASIA
4	33	JZI	Airlines	No	6.30	Online	53	18.00	Bronze Plan	ASIA

The data has read successfully,

The shape of the dataset is (3000, 10)

Info function clearly indicates the dataset has object, integer and float so we have to change the object data type to numeric value.

```
df_insured.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              3000 non-null   int64
1   Agency_Code      3000 non-null   object
2   Type             3000 non-null   object
3   Claimed          3000 non-null   object
4   Commision        3000 non-null   float64
5   Channel          3000 non-null   object
6   Duration         3000 non-null   int64
7   Sales            3000 non-null   float64
8   Product Name     3000 non-null   object
9   Destination      3000 non-null   object
dtypes: float64(2), int64(2), object(6)
memory usage: 234.5+ KB
```

No missing values in the dataset,

Check for missing value in any column

```
] df_insured.isnull().sum()
```

```
Age              0
Agency_Code     0
Type             0
Claimed          0
Commision        0
Channel          0
Duration         0
Sales            0
Product Name     0
Destination      0
dtype: int64
```

Summary of the dataset,

Summary of the data

```
]: df_insured.describe(include="all").T
```

```
at[9]:
```

	count	unique	top	freq	mean	std	min	25%	50%	75%	max
Age	3000	NaN	NaN	NaN	38.091	10.4635	8	32	36	42	84
Agency_Code	3000	4	EPX	1365	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Type	3000	2	Travel Agency	1837	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Claimed	3000	2	No	2076	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Commision	3000	NaN	NaN	NaN	14.5292	25.4815	0	0	4.63	17.235	210.21
Channel	3000	2	Online	2954	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Duration	3000	NaN	NaN	NaN	70.0013	134.053	-1	11	26.5	63	4580
Sales	3000	NaN	NaN	NaN	60.2499	70.734	0	20	33	69	539
Product Name	3000	5	Customised Plan	1136	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Destination	3000	3	ASIA	2465	NaN	NaN	NaN	NaN	NaN	NaN	NaN

We have 4 numeric values and 6 categorical values,

Agency code EPX has a frequency of 1365,

The most preferred type seems to be travel agency

Channel is online

Customized plan is the most sought plan by customers

Destination ASIA seems to be most sought destination place by customers.

We will further look at the distribution of dataset in univariate and bivariate analysis

Checking for duplicates in the dataset,

Check for duplicate data

```
: dups = df_insured.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))
```

```
Number of duplicate rows = 139
```

Removing Duplicates

since i don't find any unique identifier in the dataset to remove these duplicates these duplicates can be different customers so i'm not dropping these duplicates.


Checking for Outliers

As there is no unique identifier I'm not dropping the duplicates it may be different customer's data.

Checking for Outliers

```
[11]: plt.figure(figsize=(10,5))
      df_insured[['Age','Commision', 'Duration', 'Sales']].boxplot()
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x22d338cce88>
```



Outliers exist in almost all the numeric values.

We can treat outliers in random forest classification.

Geting unique counts of all Nominal Variables

```
]: for column in df_insured[['Agency_Code', 'Type', 'Claimed', 'Channel',
                             'Product Name', 'Destination']]:
    print(column.upper(),': ',df_insured[column].nunique())
    print(df_insured[column].value_counts().sort_values())
    print('\n')
```

AGENCY_CODE: 4

JZI 239

CWT 472

C2B 924

EPX 1365

TYPE: 2

Airlines 1163

Travel Agency 1837

CLAIMED: 2



Yes 924
No 2076

CHANNEL: 2
Offline 46
Online 2954

PRODUCT NAME: 5
Gold Plan 109
Silver Plan 427
Bronze Plan 650
Cancellation Plan 678
Customised Plan 1136

DESTINATION: 3
EUROPE 215
Americas 320
ASIA 2465

Univariate / Bivariate analysis




The box plot of the age variable shows outliers.

Spending is positively skewed - 1.149713

The dist plot shows the distribution of data from 20 to 80


In the range of 30 to 40 is where the majority of the distribution lies.



The box plot of the commision variable shows outliers.

Spending is positively skewed - 3.148858


The dist plot shows the distribution of data from 0 to 30



The box plot of the duration variable shows outliers.

Spending is positively skewed - 13.784681

The dist plot shows the distribution of data from 0 to 100



The box plot of the sales variable shows outliers.

Spending is positively skewed - 2.381148


The dist plot shows the distribution of data from 0 to 300

Categorical Variables

Agency Code


```
sns.countplot(df_insured['Agency_Code'])
```

```
]: <matplotlib.axes._subplots.AxesSubplot at 0x22d34aa6448>
```



The distribution of the agency code, shows us EPX with maximum frequency


```
sns.boxplot(data = df_insured, x='Agency_Code',y='Sales', hue='Claimed')
<matplotlib.axes._subplots.AxesSubplot at 0x22d34b391c8>
```



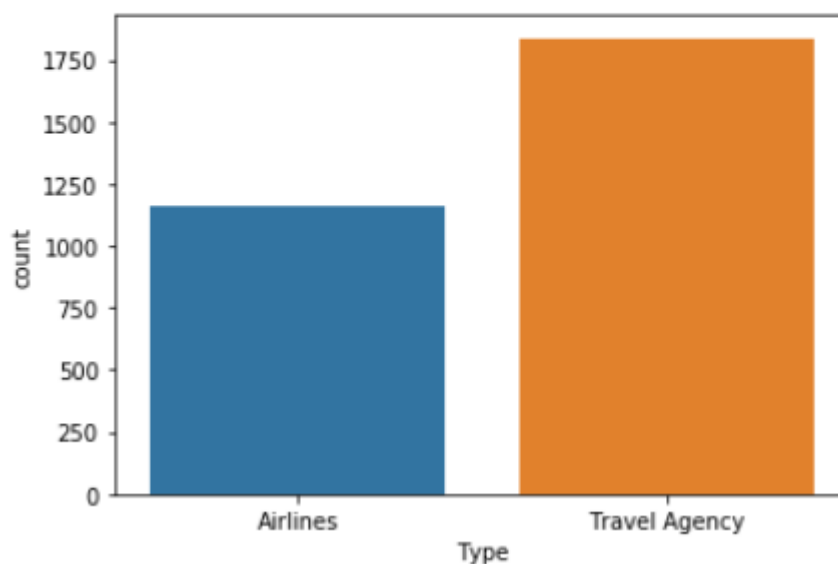
The box plot shows the split of sales with different agency code and also hue having claimed column.

It seems that C2B have claimed more claims than other agency.

Type


```
7] sns.countplot(data = df_insured, x = 'Type')
```

7]: <matplotlib.axes._subplots.AxesSubplot at 0x22d3471c788>



```
sns.boxplot(data = df_insured, x='Type',y='Sales', hue='Claimed')
```

```
]: <matplotlib.axes._subplots.AxesSubplot at 0x22d34763288>
```




The box plot shows the split of sales with different type and also hue having claimed column. We could understand airlines type has more claims.

Channel

```
sns.countplot(data = df_insured, x = 'Channel')
```

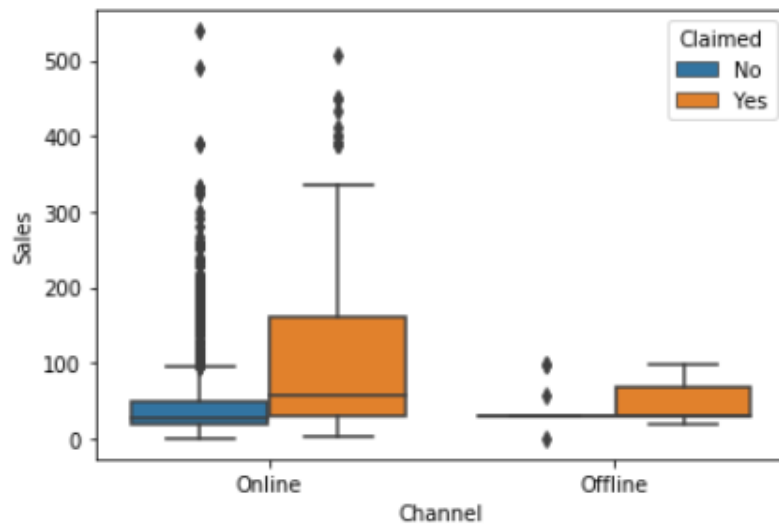
```
.9]: <matplotlib.axes._subplots.AxesSubplot at 0x22d34827e48>
```



The majority of customers have used online medium, very less with offline medium

```
1] sns.boxplot(data = df_insured, x='Channel',y='Sales', hue='Claimed')
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x22d34883588>
```




The box plot shows the split of sales with different channel and also hue having claimed column.

Product Name

```
1] sns.countplot(data = df_insured, x = 'Product Name')
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x22d348ce308>
```

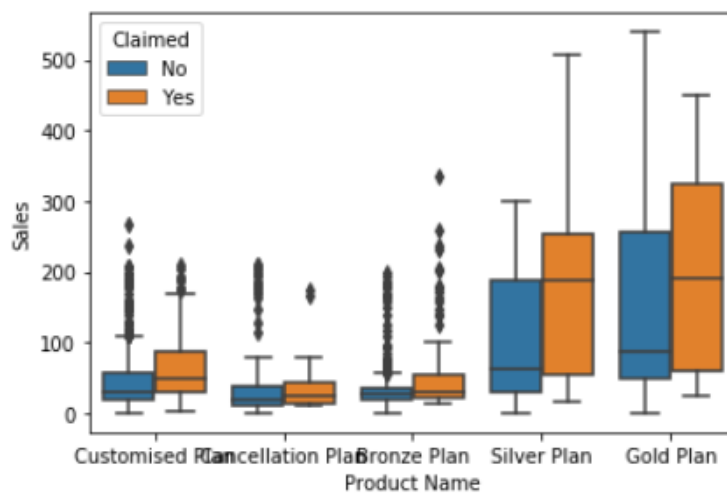


Customized plan seems to be most liked plan by customers when compared to all other plans.

```

sns.boxplot(data = df_insured, x='Product Name',y='Sales', hue='Claimed')
]: <matplotlib.axes._subplots.AxesSubplot at 0x22d349a2408>

```




The box plot shows the split of sales with different product name and also hue having claimed column.

Destination

```

sns.countplot(data = df_insured, x = 'Destination')
!3]: <matplotlib.axes._subplots.AxesSubplot at 0x22d35eba388>


```



Asia is where customers choose when compared with other destination places.

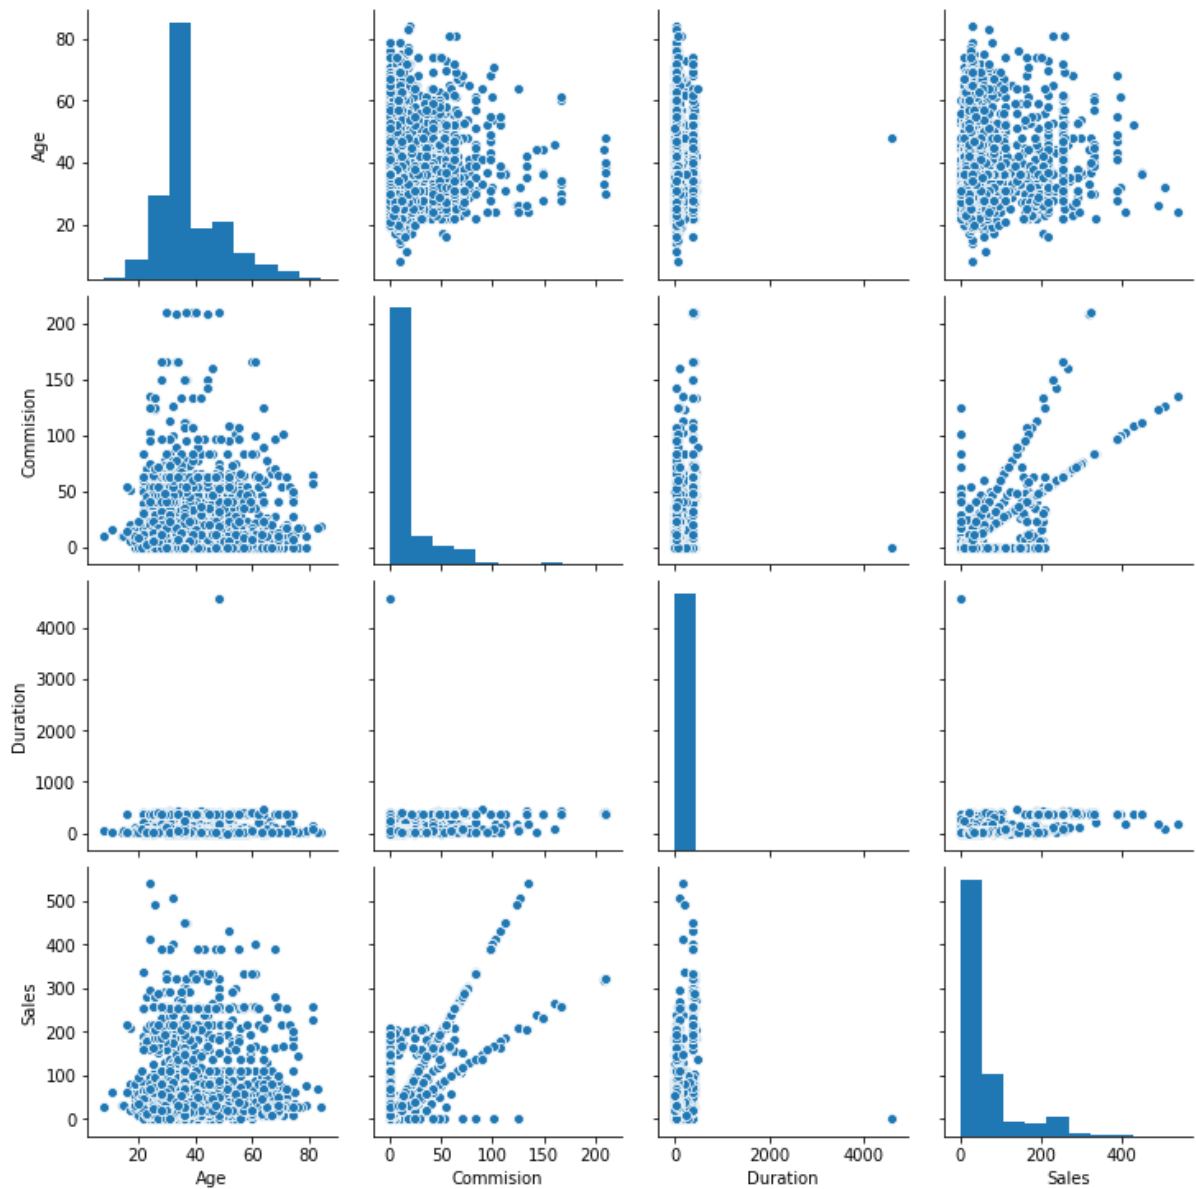
```
sns.boxplot(data = df_insured, x='Destination',y='Sales', hue='Claimed')
```

```
Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x22d35ef1ec8>
```




The box plot shows the split of sales with different destination and also hue having claimed column.

Checking pairwise distribution of the continuous variables



Checking for Correlations



Not much of multi collinearity observed

No negative correlation

Only positive correlation

Converting all objects to categorical codes

```

> for feature in df_insured.columns:
    if df_insured[feature].dtype == 'object':
        print('\n')
        print('feature:', feature)
        print(pd.Categorical(df_insured[feature].unique()))
        print(pd.Categorical(df_insured[feature].unique()).codes)
        df_insured[feature] = pd.Categorical(df_insured[feature]).codes
  
```

To build our models we are changing the object data type to numeric values.

feature: Agency Code

[C2B, EPX, CWT, JZI]

Categories (4, object): [C2B, CWT, EPX, JZI]

[0 2 1 3]

Feature: Type

[Airlines, Travel Agency]

Categories (2, object): [Airlines, Travel Agency]

[0 1]

Feature: Claimed

[No, Yes]

Categories (2, object): [No, Yes]

[0 1]

Feature: Channel

[Online, Offline]

Categories (2, object): [Offline, Online]

[1 0]

Feature: Product Name

[Customised Plan, Cancellation Plan, Bronze Plan, Silver Plan, Gold Plan]

Categories (5, object): [Bronze Plan, Cancellation Plan, Customised Plan, Gold Plan, Silver Plan]

[2 1 0 4 3]

Feature: Destination

[ASIA, Americas, EUROPE]

Categories (3, object): [ASIA, Americas, EUROPE]

[0 1 2]

Checking the info

```
df_insured.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              3000 non-null   int64
1   Agency_Code      3000 non-null   int8
2   Type             3000 non-null   int8
3   Claimed          3000 non-null   int8
4   Commision        3000 non-null   float64
5   Channel          3000 non-null   int8
6   Duration         3000 non-null   int64
7   Sales            3000 non-null   float64
8   Product Name     3000 non-null   int8
9   Destination      3000 non-null   int8
dtypes: float64(2), int64(2), int8(6)
memory usage: 111.5 KB
```

```
df_insured.head()
```

0]:

	Age	Agency_Code	Type	Claimed	Commision	Channel	Duration	Sales	Product Name	Destination
0	48	0	0	0	0.70	1	7	2.51	2	0
1	36	2	1	0	0.00	1	34	20.00	2	0
2	39	1	1	0	5.94	1	3	9.90	2	1
3	36	2	1	0	0.00	1	4	26.00	1	0
4	33	3	0	0	6.30	1	53	18.00	0	0

Proportion of 1s and 0s

```
df_insured.Claimed.value_counts(normalize=True)
```

```
30]: 0    0.692
      1    0.308
      Name: Claimed, dtype: float64
```

Checking the proportion of 1s and 2s in the dataset. That is our target column.

2.2 Data Split: Split the data into test and train, build classification model CART, Random Forest, Artificial Neural Network

Extracting the target column into separate vectors for training set and test set

```
X = df_insured.drop("Claimed", axis=1)
y = df_insured.pop("Claimed")
X.head()
```

```
31]:
```

	Age	Agency_Code	Type	Commision	Channel	Duration	Sales	Product Name	Destination
0	48	0	0	0.70	1	7	2.51	2	0
1	36	2	1	0.00	1	34	20.00	2	0
2	39	1	1	5.94	1	3	9.90	2	1
3	36	2	1	0.00	1	4	26.00	1	0
4	33	3	0	6.30	1	53	18.00	0	0

For training and testing purpose we are splitting the dataset into train and test data in the ratio 70:30.

Splitting data into training and test set

```
X_train, X_test, train_labels, test_labels = train_test_split(X, y, test_size=.30, random_state=1)
```

Checking the dimensions of the training and test data

```

: ▶ print('X_train',X_train.shape)
   print('X_test',X_test.shape)
   print('train_labels',train_labels.shape)
   print('test_labels',test_labels.shape)

X_train (2100, 9)
X_test (900, 9)
train_labels (2100,)
test_labels (900,)

```

We have bifurcated the dataset into train and test.

We have also taken out the target column out of train and test data into separate vector for evaluation purposes.

MODEL 1

Building a Decision Tree Classifier

```

▶ dt_model = DecisionTreeClassifier(criterion = 'gini' )

```

```

▶ dt_model.fit(X_train, train_labels)

```

```

35]: DecisionTreeClassifier()

```

CHECKING THE FEATURE

```

▶ print (pd.DataFrame(dt_model.feature_importances_, columns = ["Imp"],
                    index = X_train.columns).sort_values('Imp',ascending=False))

```

	Imp
Duration	0.276811
Agency_Code	0.194356
Sales	0.194228
Age	0.163714
Commision	0.102841
Product Name	0.038334
Destination	0.019359
Channel	0.007262
Type	0.003095

OPTIMAL VALUES FOR DECISSION TREE,

GRID SEARCH FOR FINDING,

Grid Search for finding out the optimal values for the hyper parameters

```
from sklearn.model_selection import GridSearchCV

param_grid = {
    'max_depth': [4, 5, 6],
    'min_samples_leaf': [20, 40, 60, 70],
    'min_samples_split': [150, 200, 250, 300,]
}

dt_model = DecisionTreeClassifier()

grid_search = GridSearchCV(estimator = dt_model, param_grid = param_grid, cv = 10)
```

FITTING THE OPTMAL VALUES TO THE TRAINING DATASET

```
grid_search.fit(X_train, train_labels)

: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(),
               param_grid={'max_depth': [4, 5, 6],
                           'min_samples_leaf': [20, 40, 60, 70],
                           'min_samples_split': [150, 200, 250, 300]})
```

BEST GRID

```
best_grid

DecisionTreeClassifier(max_depth=4, min_samples_leaf=20, min_samples_split=150)
```

Regularising the Decision Tree

Adding Tuning Parameters

```
reg_dt_model = DecisionTreeClassifier(criterion = 'gini', max_depth = 4,min_samples_leaf=20,min_samples_split=150)
```

```
reg_dt_model.fit(X_train, train_labels)
```

```
DecisionTreeClassifier(max_depth=4, min_samples_leaf=20, min_samples_split=150)
```

Generating New Tree

```
insurance_prediction_tree_regularized = open('C:\\Users\\WELCOME\\Downloads\\PYTHON FILES\\4.Data Mining\\Project\\insurance_p
dot_data = tree.export_graphviz(reg_dt_model, out_file= insurance_prediction_tree_regularized , feature_names = list(X_train)

insurance_prediction_tree_regularized.close()
dot_data
```

Variable Importance

```
: ▶ print (pd.DataFrame(reg_dt_model.feature_importances_, columns = ["Imp"],
                        index = X_train.columns).sort_values('Imp',ascending=False))
```

```

                Imp
Agency_Code    0.616392
Sales           0.252286
Product Name    0.077771
Commision       0.022912
Duration        0.022624
Age             0.008015
Type            0.000000
Channel         0.000000
Destination     0.000000
```

Predicting on Training dataset for Decission Tree

```
▶ ytrain_predict_dt = reg_dt_model.predict(X_train)
```

```
▶ ytest_predict_dt = reg_dt_model.predict(X_test)
```

MODEL 2

Building a Ensemble RandomForest Classifier

```
▶ df_insured_rf=df_original.copy()
df_insured_rf.head()
```

9]:

	Age	Agency_Code	Type	Claimed	Commision	Channel	Duration	Sales	Product Name	Destination
0	48	C2B	Airlines	No	0.70	Online	7	2.51	Customised Plan	ASIA
1	36	EPX	Travel Agency	No	0.00	Online	34	20.00	Customised Plan	ASIA
2	39	CWT	Travel Agency	No	5.94	Online	3	9.90	Customised Plan	Americas
3	36	EPX	Travel Agency	No	0.00	Online	4	26.00	Cancellation Plan	ASIA
4	33	JZI	Airlines	No	6.30	Online	53	18.00	Bronze Plan	ASIA


TREATING OUTLIERS FOR RANDOM FOREST

```

▶ def treat_outlier(col):
    sorted(col)
    Q1,Q3=np.percentile(col,[25,75])
    IQR=Q3-Q1
    lower_range= Q1-(1.5 * IQR)
    upper_range= Q3+(1.5 * IQR)
    return lower_range, upper_range

▶ for feature in df_insured_rf[['Age','Commision', 'Duration', 'Sales']]:
    lr,ur=treat_outlier(df_insured_rf[feature])
    df_insured_rf[feature]=np.where(df_insured_rf[feature]>ur,ur,df_insured_rf[feature])
    df_insured_rf[feature]=np.where(df_insured_rf[feature]<lr,lr,df_insured_rf[feature])
```

BOX PLOT TO CHECK PRESENCE OF OUTLIERS



RANDOM FOREST CLASSIFIER

```
X_train, X_test, train_labels, test_labels = train_test_split(X_rf, y_rf, test_size=.30, random_state=1)
```

```
rfcl = RandomForestClassifier(n_estimators = 100,max_features=6,random_state=1)
rfcl = rfcl.fit(X_train, train_labels)
```

```
rfcl
```

TO FIND OPTIMAL NUMBERS USING GRID SEARCH

Grid Search for finding out the optimal values for the hyper parameters

```
param_grid_rfcl = {
    'max_depth': [6],##20,30,40
    'max_features': [4],## 7,8,9
    'min_samples_leaf': [8],## 50,100
    'min_samples_split': [45], ## 60,70
    'n_estimators': [100] ## 100,200
}

rfcl = RandomForestClassifier(random_state=1)

grid_search_rfcl = GridSearchCV(estimator = rfcl, param_grid = param_grid_rfcl, cv = 10)
```

FITTING THE MODEL TO RFCL VALUES OBTAINED BY OPTIMAL GRID SEARCH METHOD

```
grid_search_rfcl.fit(X_train, train_labels)

GridSearchCV(cv=10, estimator=RandomForestClassifier(random_state=1),
             param_grid={'max_depth': [6], 'max_features': [4],
                          'min_samples_leaf': [8], 'min_samples_split': [45],
                          'n_estimators': [100]})
```

BEST GRID VALUES

```
best_grid_rf
```

```
] RandomForestClassifier(max_depth=6, max_features=4, min_samples_leaf=8,
                        min_samples_split=45, random_state=1)
```

Predicting on Training dataset for Random Forest

```
|: best_grid_rf.predict(X_train)
```

```
|: best_grid_rf.predict(X_test)
```

MODEL 3

Building a Neural Network Classifier

BEFORE BUILDING THE MODEL

WE SCALE THE VALUES, TO STANDARD SCALE USING MINMAXSCALER

```
sc = StandardScaler()
```

```
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

AFTER SCALING WE ARE TRANSFORMING THE SAME TO THE TEST DATA

```
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

MLP CLASSIFIER

```
clf = MLPClassifier(hidden_layer_sizes=100, max_iter=5000,
                    solver='sgd', verbose=True, random_state=21, tol=0.01)
```

TRAINING THE MODEL

```

| clf.fit(X_trains, train_labels)

Iteration 1, loss = 0.64244509
Iteration 2, loss = 0.62392631
Iteration 3, loss = 0.60292414
Iteration 4, loss = 0.58458220
Iteration 5, loss = 0.56914550
Iteration 6, loss = 0.55651481
Iteration 7, loss = 0.54598011
Iteration 8, loss = 0.53752961
Iteration 9, loss = 0.53051147
Iteration 10, loss = 0.52440802
Iteration 11, loss = 0.51934384
Iteration 12, loss = 0.51483466
Iteration 13, loss = 0.51108343
Iteration 14, loss = 0.50763356
Iteration 15, loss = 0.50476577
Iteration 16, loss = 0.50218466
Iteration 17, loss = 0.49989583
Iteration 18, loss = 0.49786338
Training loss did not improve more than tol=0.010000 for 10 consecutive epochs. Stopping.

```

GRID SEARCH

Grid Search for finding out the optimal values for the hyper parameters

```

:  | param_grid = {
      'hidden_layer_sizes': [200], # 50, 200
      'max_iter': [2500], #5000,2500
      'solver': ['adam'], #sgd
      'tol': [0.01],
    }

    nncl = MLPClassifier(random_state=1)

    grid_search = GridSearchCV(estimator = nncl, param_grid = param_grid, cv = 10)

```

FITTING THE MODEL USING THE OPTIMAL VALUES FROM GRID SEARCH

```

:  | grid_search.fit(X_trains, train_labels)

:  GridSearchCV(cv=10, estimator=MLPClassifier(random_state=1),
      param_grid={'hidden_layer_sizes': [200], 'max_iter': [2500],
        'solver': ['adam'], 'tol': [0.01]})

```

BEST GRID VALUES,

```

:  | best_grid_ann= grid_search.best_estimator_
      best_grid_ann

:  MLPClassifier(hidden_layer_sizes=200, max_iter=2500, random_state=1, tol=0.01)

```

Predicting on Training dataset for Neural Network Classifier

```

ytrain_predict = best_grid_ann.predict(X_trains)
ytest_predict = best_grid_ann.predict(X_tests)

```

2.3 Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC_AUC score for each model

DECISION TREE PREDICTION

Predicting on Training dataset for Decision Tree

```

ytrain_predict_dt = reg_dt_model.predict(X_train)

print('ytrain_predict',ytrain_predict_dt.shape)

ytrain_predict (2100,)

```

ACCURACY

```

cart_train_acc = reg_dt_model.score(X_train,train_labels)
cart_train_acc

```

```
4]: 0.7933333333333333
```

CONFUSION MATRIX

```

print(classification_report(train_labels, ytrain_predict_dt))

```

	precision	recall	f1-score	support
0	0.84	0.87	0.85	1471
1	0.67	0.62	0.64	629
accuracy			0.79	2100
macro avg	0.75	0.74	0.75	2100
weighted avg	0.79	0.79	0.79	2100

```

confusion_matrix(train_labels, ytrain_predict_dt)

```

```
3]: array([[1275, 196],
          [ 238, 391]], dtype=int64)
```


Model Evaluation for Decision Tree

AUC and ROC for the training data for Decision Tree


```

# predict probabilities
probs = reg_dt_model.predict_proba(X_train)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
cart_train_auc = roc_auc_score(train_labels, probs)
print('AUC: %.3f' % cart_train_auc)
# calculate roc curve
cart_train_fpr, cart_train_tpr, cart_train_thresholds = roc_curve(train_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(cart_train_fpr, cart_train_tpr)

```

AUC: 0.827

[<matplotlib.lines.Line2D at 0x22d36d29fc8>]



```

cart_metrics=classification_report(train_labels, ytrain_predict_dt,output_dict=True)
df=pd.DataFrame(cart_metrics).transpose()
cart_train_f1=round(df.loc["1"][2],2)
cart_train_recall=round(df.loc["1"][1],2)
cart_train_precision=round(df.loc["1"][0],2)
print ('cart_train_precision ',cart_train_precision)
print ('cart_train_recall ',cart_train_recall)
print ('cart_train_f1 ',cart_train_f1)

```

cart_train_precision 0.67
 cart_train_recall 0.62
 cart_train_f1 0.64

AUC and ROC for the test data for Decision Tree

```

▶ # predict probabilities
probs = reg_dt_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
cart_test_auc = roc_auc_score(test_labels, probs)
print('AUC: %.3f' % cart_test_auc)
# calculate roc curve
cart_test_fpr, cart_test_tpr, cart_testthresholds = roc_curve(test_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(cart_test_fpr, cart_test_tpr)

```

AUC: 0.790

```

▶ cart_metrics=classification_report(test_labels, ytest_predict_dt,output_dict=True)
df=pd.DataFrame(cart_metrics).transpose()
cart_test_precision=round(df.loc["1"][0],2)
cart_test_recall=round(df.loc["1"][1],2)
cart_test_f1=round(df.loc["1"][2],2)
print ('cart_test_precision ',cart_test_precision)
print ('cart_test_recall ',cart_test_recall)
print ('cart_test_f1 ',cart_test_f1)

```

```

cart_test_precision 0.71
cart_test_recall    0.53
cart_test_f1        0.6

```

MODEL 2 PREDICTION RANDOM FOREST

Predicting on Training dataset for Random Forest

```

▶ ytrain_predict_rf = best_grid_rf.predict(X_train)

```

```

▶ print('ytrain_predict',ytrain_predict_rf.shape)

```

```

ytrain_predict (2100,)

```

ACCURACY

```

▶ rf_train_acc = best_grid_rf.score(X_train,train_labels)
rf_train_acc

```

```

]: 0.8123809523809524

```

CONFUSION MATRIX

```
print(classification_report(train_labels, ytrain_predict_rf))
```

	precision	recall	f1-score	support
0	0.84	0.90	0.87	1471
1	0.73	0.60	0.66	629
accuracy			0.81	2100
macro avg	0.78	0.75	0.76	2100
weighted avg	0.81	0.81	0.81	2100


```
confusion_matrix(train_labels, ytrain_predict_rf)
```

```
array([[1331, 140],
       [ 254, 375]], dtype=int64)
```

Model Evaluation for Random Forest

AUC and ROC for the training data for Random Forest

```
# predict probabilities
probs = best_grid_rf.predict_proba(X_train)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
rf_train_auc = roc_auc_score(train_labels, probs)
print('AUC: %.3f' % rf_train_auc)
# calculate roc curve
rf_train_fpr, rf_train_tpr, rf_train_thresholds = roc_curve(train_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(rf_train_fpr, rf_train_tpr, color='green')
```



```

rf_metrics=classification_report(train_labels, ytrain_predict_rf,output_dict=True)
df=pd.DataFrame(rf_metrics).transpose()
rf_train_f1=round(df.loc["1"][2],2)
rf_train_recall=round(df.loc["1"][1],2)
rf_train_precision=round(df.loc["1"][0],2)
print ('rf_train_precision ',rf_train_precision)
print ('rf_train_recall ',rf_train_recall)
print ('rf_train_f1 ',rf_train_f1)

rf_train_precision 0.73
rf_train_recall 0.6
rf_train_f1 0.66

```

Predicting on Test dataset for Random Forest

```

: ▶ ytest_predict_rf = best_grid_rf.predict(X_test)

: ▶ print('ytest_predict_rf',ytest_predict_rf.shape)

ytest_predict_rf (900,)

```

ACCURACY

```

rf_test_acc = best_grid_rf.score(X_test,test_labels)
rf_test_acc

0.7733333333333333

```

CONFUSION MATRIX

```
print(classification_report(test_labels, ytest_predict_rf))
```

	precision	recall	f1-score	support
0	0.79	0.91	0.84	605
1	0.73	0.49	0.59	295
accuracy			0.77	900
macro avg	0.76	0.70	0.71	900
weighted avg	0.77	0.77	0.76	900


AUC and ROC for the test data for Random Forest

```

# predict probabilities
probs = best_grid_rf.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
rf_test_auc = roc_auc_score(test_labels, probs)
print('AUC: %.3f' % rf_test_auc)
# calculate roc curve
rf_test_fpr, rf_test_tpr, rf_testthresholds = roc_curve(test_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(rf_test_fpr, rf_test_tpr, color='green')

```

AUC: 0.818



```

rf_metrics=classification_report(test_labels, ytest_predict_rf,output_dict=True)
df=pd.DataFrame(rf_metrics).transpose()
rf_test_precision=round(df.loc["1"][0],2)
rf_test_recall=round(df.loc["1"][1],2)
rf_test_f1=round(df.loc["1"][2],2)
print ('rf_test_precision ',rf_test_precision)
print ('rf_test_recall ',rf_test_recall)
print ('rf_test_f1 ',rf_test_f1)

```

rf_test_precision 0.73

rf_test_recall 0.49

rf_test_f1 0.59

MODEL 3

ANN

Predicting on Training dataset for Neural Network Classifier

```

ytrain_predict_ann = best_grid_ann.predict(X_trains)

```

```

print('ytrain_predict',ytrain_predict_ann.shape)

```

ytrain_predict (2100,)

CONFUSION MATRIX

```
confusion_matrix(train_labels,ytrain_predict_ann)
```

```
]: array([[1317, 154],
        [ 303, 326]], dtype=int64)
```

ACCURACY

```
ann_train_acc=best_grid_ann.score(X_trains,train_labels)
ann_train_acc
```

```
: 0.7823809523809524
```

```
print(classification_report(train_labels,ytrain_predict_ann))
```


	precision	recall	f1-score	support
0	0.81	0.90	0.85	1471
1	0.68	0.52	0.59	629
accuracy			0.78	2100
macro avg	0.75	0.71	0.72	2100
weighted avg	0.77	0.78	0.77	2100

Model Evaluation for Neural Network Classifier

AUC and ROC for the training data for Neural Network Classifier

```
# predict probabilities
probs = best_grid_ann.predict_proba(X_trains)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
ann_train_auc = roc_auc_score(train_labels, probs)
print('AUC: %.3f' % ann_train_auc)
# calculate roc curve
ann_train_fpr, ann_train_tpr, ann_train_thresholds = roc_curve(train_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(ann_train_fpr, ann_train_tpr, color='yellow')
```

AUC: 0.823



Predicting on Test dataset for Neural Network Classifier

```
▶ ytest_predict_ann = best_grid_ann.predict(X_tests)
```

```
▶ print('ytest_predict_ann', ytest_predict_ann.shape)
```

```
ytest_predict_ann (900,)
```

ACCURACY

```
| ann_test_acc = best_grid_ann.score(X_tests, test_labels)
| ann_test_acc
```

```
0.7622222222222222
```

CONFUSION MATRIX


```
▶ confusion_matrix(test_labels, ytest_predict_ann)
```

```
}: array([[557, 48],
         [166, 129]], dtype=int64)
```

AUC and ROC for the test data for Neural Network Classifier



```
▶ # predict probabilities
  probs = best_grid_ann.predict_proba(X_tests)
  # keep probabilities for the positive outcome only
  probs = probs[:, 1]
  # calculate AUC
  ann_test_auc = roc_auc_score(test_labels, probs)
  print('AUC: %.3f' % ann_test_auc)
  # calculate roc curve
  ann_test_fpr, ann_test_tpr, ann_test_thresholds = roc_curve(test_labels, probs)
  plt.plot([0, 1], [0, 1], linestyle='--')
  # plot the roc curve for the model
  plt.plot(ann_test_fpr, ann_test_tpr, color='yellow')
```

```
AUC: 0.806
```



2.4 Final Model: Compare all the model and write an inference which model is best/optimized.

	CART Train	CART Test	Random Forest Train	Random Forest Test	Neural Network Train	Neural Network Test
Accuracy	0.79	0.77	0.81	0.77	0.78	0.76
AUC	0.83	0.79	0.86	0.82	0.82	0.81
Recall	0.62	0.53	0.60	0.49	0.52	0.44
Precision	0.67	0.71	0.73	0.73	0.68	0.73
F1 Score	0.64	0.60	0.66	0.59	0.59	0.55



CONCLUSION:

I am selecting the RF model, as it has better accuracy, precision, recall, and f1 score better than other two CART & NN.

2.5 Inference: Based on the whole Analysis, what are the business insights and recommendations?

Looking at the model, more data will help us understand and predict models better.

Streamlining online experiences benefitted customers, leading to an increase in conversions, which subsequently raised profits.

- As per the data 90% of insurance is done by online channel.
- Other interesting fact, is almost all the offline business has a claimed associated
- Need to train the JZI agency resources to pick up sales as they are in bottom, need to run promotional marketing campaign or evaluate if we need to tie up with alternate agency
- Also based on the model we are getting 80%accuracy, so we need customer books airline tickets or plans, cross sell the insurance based on the claim data pattern.
- Other interesting fact is more sales happen via Agency than Airlines and the trend shows the claim are processed more at Airline. So we may need to deep dive into the process to understand the workflow and why?

Key performance indicators (KPI) The KPI's of insurance claims are

- Increase customer satisfaction which in fact will give more revenue
- Combat fraud transactions, deploy measures to avoid fraudulent transactions at earliest
- Optimize claims recovery method
- Reduce claim handling costs.